


Neural Networks



Aalto-yliopisto
Aalto-universitetet
Aalto University


Medical Image Analysis
Koen Van Leemput
Fall 2023

Focus: automatic segmentation




Remember linear regression?

- Let $\mathbf{x} = (x_1, \dots, x_D)^T$ denote an input vector in a D -dimensional space
- Given N measurements $\{t_n\}_{n=1}^N$ at inputs $\{\mathbf{x}_n\}_{n=1}^N$, what is t at a new input \mathbf{x} ?







$$y(\mathbf{x}, \mathbf{w}) = \sum_{m=0}^{M-1} w_m \phi_m(\mathbf{x})$$

tunable weights







Remember the Gaussian mixture model?



Posterior using Bayes' rule:
$$p(l = k|d, \theta) = \frac{\mathcal{N}(d|\mu_k, \sigma_k^2)\pi_k}{\sum_{k'} \mathcal{N}(d|\mu_{k'}, \sigma_{k'}^2)\pi_{k'}}$$

Remember the Gaussian mixture model?






new notation/terminology:

- “training samples”
- “ $t=1$ ” if $l=1$
- “ $t=0$ ” if $l=2$


Posterior using Bayes' rule:
$$p(l = k|d, \theta) = \frac{\mathcal{N}(d|\mu_k, \sigma_k^2)\pi_k}{\sum_{k'} \mathcal{N}(d|\mu_{k'}, \sigma_{k'}^2)\pi_{k'}}$$

Remember the Gaussian mixture model?



This lecture: can we get a “classifier”

$p(l|d, \theta)$
directly
without a model?




new notation/terminology:

- “training samples”
- “t=1” if $l=1$
- “t=0” if $l=2$

Posterior using Bayes’ rule:
$$p(l = k|d, \theta) = \frac{\mathcal{N}(d|\mu_k, \sigma_k^2)\pi_k}{\sum_{k'} \mathcal{N}(d|\mu_{k'}, \sigma_{k'}^2)\pi_{k'}}$$


Logistic regression

- Logistic function as a “squashing” function



$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$




Logistic regression



- $p(t = 1|\mathbf{x}, \boldsymbol{\theta}) = f(\mathbf{x})$ where $f(\mathbf{x}) = \sigma \left(\sum_{m=0}^{M-1} w_m \phi_m(\mathbf{x}) \right)$
- Of course: $p(t = 0|\mathbf{x}, \boldsymbol{\theta}) = 1 - p(t = 1|\mathbf{x}, \boldsymbol{\theta}) = 1 - f(\mathbf{x})$



Voxel-based classifier

- Training data $\{\mathbf{x}_n, t_n\}_{n=1}^N$ with $\mathbf{x}_n = d_n$ (i.e., $D = 1$) and $t_n \in \{0, 1\}$
- Estimate parameters $\boldsymbol{\theta} = (w_0, \dots, w_{M-1})^T$ by maximizing the likelihood $\prod_{n=1}^N p(t_n | \mathbf{x}_n, \boldsymbol{\theta})$





Voxel-based classifier

- Training data $\{\mathbf{x}_n, t_n\}_{n=1}^N$ with $\mathbf{x}_n = d_n$ (i.e., $D = 1$) and $t_n \in \{0, 1\}$
- Estimate parameters $\boldsymbol{\theta} = (w_0, \dots, w_{M-1})^T$ by maximizing the likelihood $\prod_{n=1}^N p(t_n | \mathbf{x}_n, \boldsymbol{\theta})$





Voxel-based classifier

- Training data $\{\mathbf{x}_n, t_n\}_{n=1}^N$ with $\mathbf{x}_n = d_n$ (i.e., $D = 1$) and $t_n \in \{0, 1\}$
- Estimate parameters $\boldsymbol{\theta} = (w_0, \dots, w_{M-1})^T$ by maximizing the likelihood $\prod_{n=1}^N p(t_n | \mathbf{x}_n, \boldsymbol{\theta})$





Voxel-based classifier

- Training data $\{\mathbf{x}_n, t_n\}_{n=1}^N$ with $\mathbf{x}_n = d_n$ (i.e., $D = 1$) and $t_n \in \{0, 1\}$
- Estimate parameters $\boldsymbol{\theta} = (w_0, \dots, w_{M-1})^T$ by maximizing the likelihood $\prod_{n=1}^N p(t_n | \mathbf{x}_n, \boldsymbol{\theta})$




Voxel-based classifier

- Once trained keep the classifier $p(l = 1|d, \hat{\theta})$
- Simply apply it to new data



$$p(l = 1|d, \hat{\theta})$$



$$p(l = 1|d, \hat{\theta}) > 0.5$$

Optimization algorithm for training

- Maximizing the likelihood function $\prod_{n=1}^N p(t_n|\mathbf{x}_n, \boldsymbol{\theta})$ is equivalent to minimizing

$$E_N(\boldsymbol{\theta}) = -\log \prod_{n=1}^N p(t_n|\mathbf{x}_n, \boldsymbol{\theta}) = -\sum_{n=1}^N \{t_n \log f(\mathbf{x}_n) + (1 - t_n) \log [1 - f(\mathbf{x}_n)]\}$$

step size (user-specified)

- Gradient descent: $\boldsymbol{\theta}^{(\tau+1)} = \boldsymbol{\theta}^{(\tau)} - \nu \nabla E_N(\boldsymbol{\theta}^{(\tau)})$ with gradient $\nabla E_N(\boldsymbol{\theta}) = \frac{\partial E_N}{\partial \boldsymbol{\theta}}$
- Stochastic gradient descent: use only $N' \ll N$ randomly sampled training points, and approximate:

$$\nabla E_N(\boldsymbol{\theta}) \simeq \frac{N}{N'} \nabla E_{N'}(\boldsymbol{\theta})$$

More fun: patch-based classifier

- Classify 3x3 image “patches”: intensity of the pixel to be classified + intensities of 8 neighboring pixels
- \mathbf{x} is now a 9-dimensional vector ($D = 9$), but otherwise everything is the same:


$$p(t = 1 | \mathbf{x}, \hat{\boldsymbol{\theta}}) = \sigma \left(\sum_{m=0}^{M-1} \hat{w}_m \phi_m(\mathbf{x}) \right)$$

- But how to choose basis functions $\phi_m(\mathbf{x})$ in a 9-dimensional space?



Basis functions in high dimensions?

- Idea: remember the “separable basis functions” trick?




Question:
does this work in 9D?



“making” sixteen 2D basis functions
out of two sets of four 1D basis functions

Basis functions in high dimensions?

- Idea: remember the “separable basis functions” trick?



“making” sixteen 2D basis functions
out of two sets of four 1D basis functions

Question:
does this work in 9D?



No!

$4^9 = 262144$
basis functions!

Adaptive basis functions

- Introduce extra parameters that alter the *form* of a limited set of basis functions
- Prototypical example:


$$\phi_m(\mathbf{x}) = \begin{cases} 1 & \text{if } m = 0, \\ \sigma \left(\sum_{d=1}^D \beta_{m,d} x_d + \beta_{m,0} \right) & \text{otherwise} \end{cases}$$

extra parameters

- All parameters ($\{\beta_{m,d}\}$ and $\{w_m\}$) are optimized together during training (stochastic gradient descent)


Adaptive basis functions (D=1)

$$\sigma(1x + 0)$$




Adaptive basis functions (D=1)

$$\sigma(0.3x + 0)$$







Adaptive basis functions (D=1)

$$\sigma(1x + 10.5)$$



Adaptive basis functions (D=2)



Feed-forward neural network


So the model is: $p(t = 1|\mathbf{x}, \boldsymbol{\theta}) = \sigma \left(\sum_{m=0}^{M-1} \tilde{w}_m \phi_m(\mathbf{x}) \right)$

parameters

with basis functions $\phi_m(\mathbf{x}) = \begin{cases} 1 & \text{if } m = 0, \\ \sigma \left(\sum_{d=1}^D \beta_{m,d} x_d + \beta_{m,0} \right) & \text{otherwise} \end{cases}$


Feed-forward neural network

Graphical representation of our 3x3 patch-based classifier ($D=9$ and $M=4$):



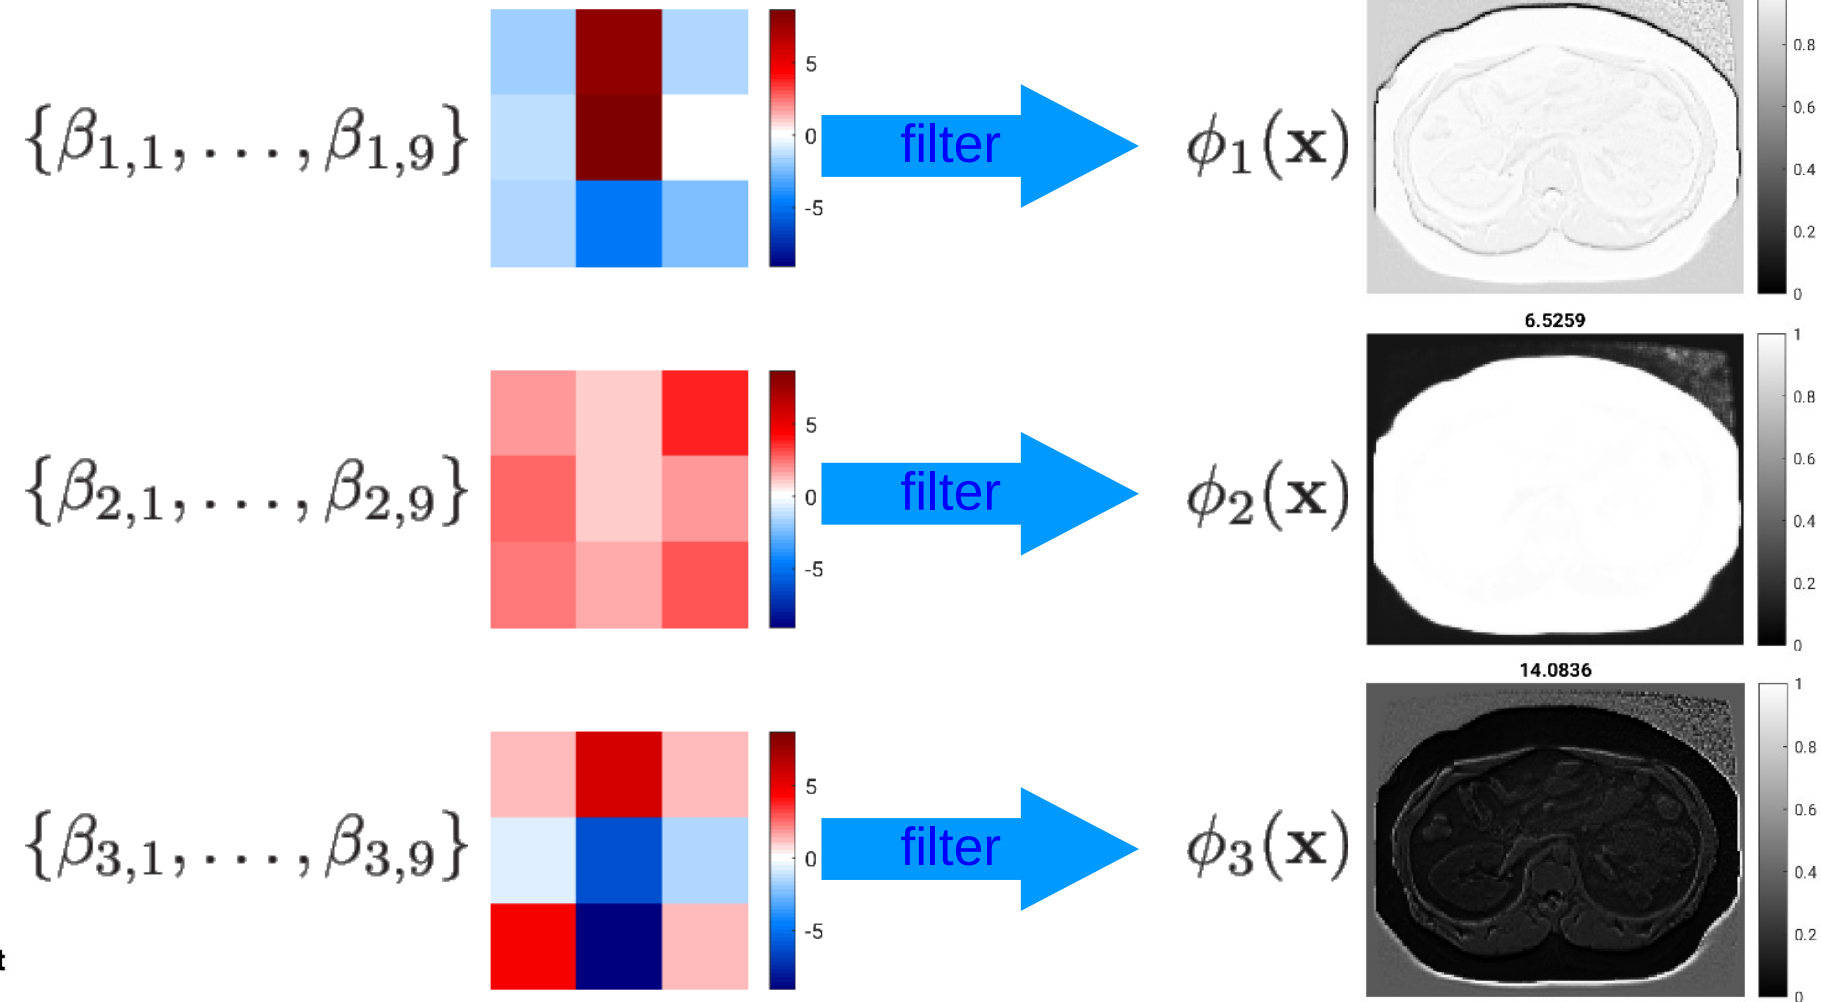

Can insert more than one "hidden" layer ("deep learning")

Applying
the trained
classifier
on new
data:




$$p(y = 1 | \mathbf{x}, \hat{\boldsymbol{\theta}})$$

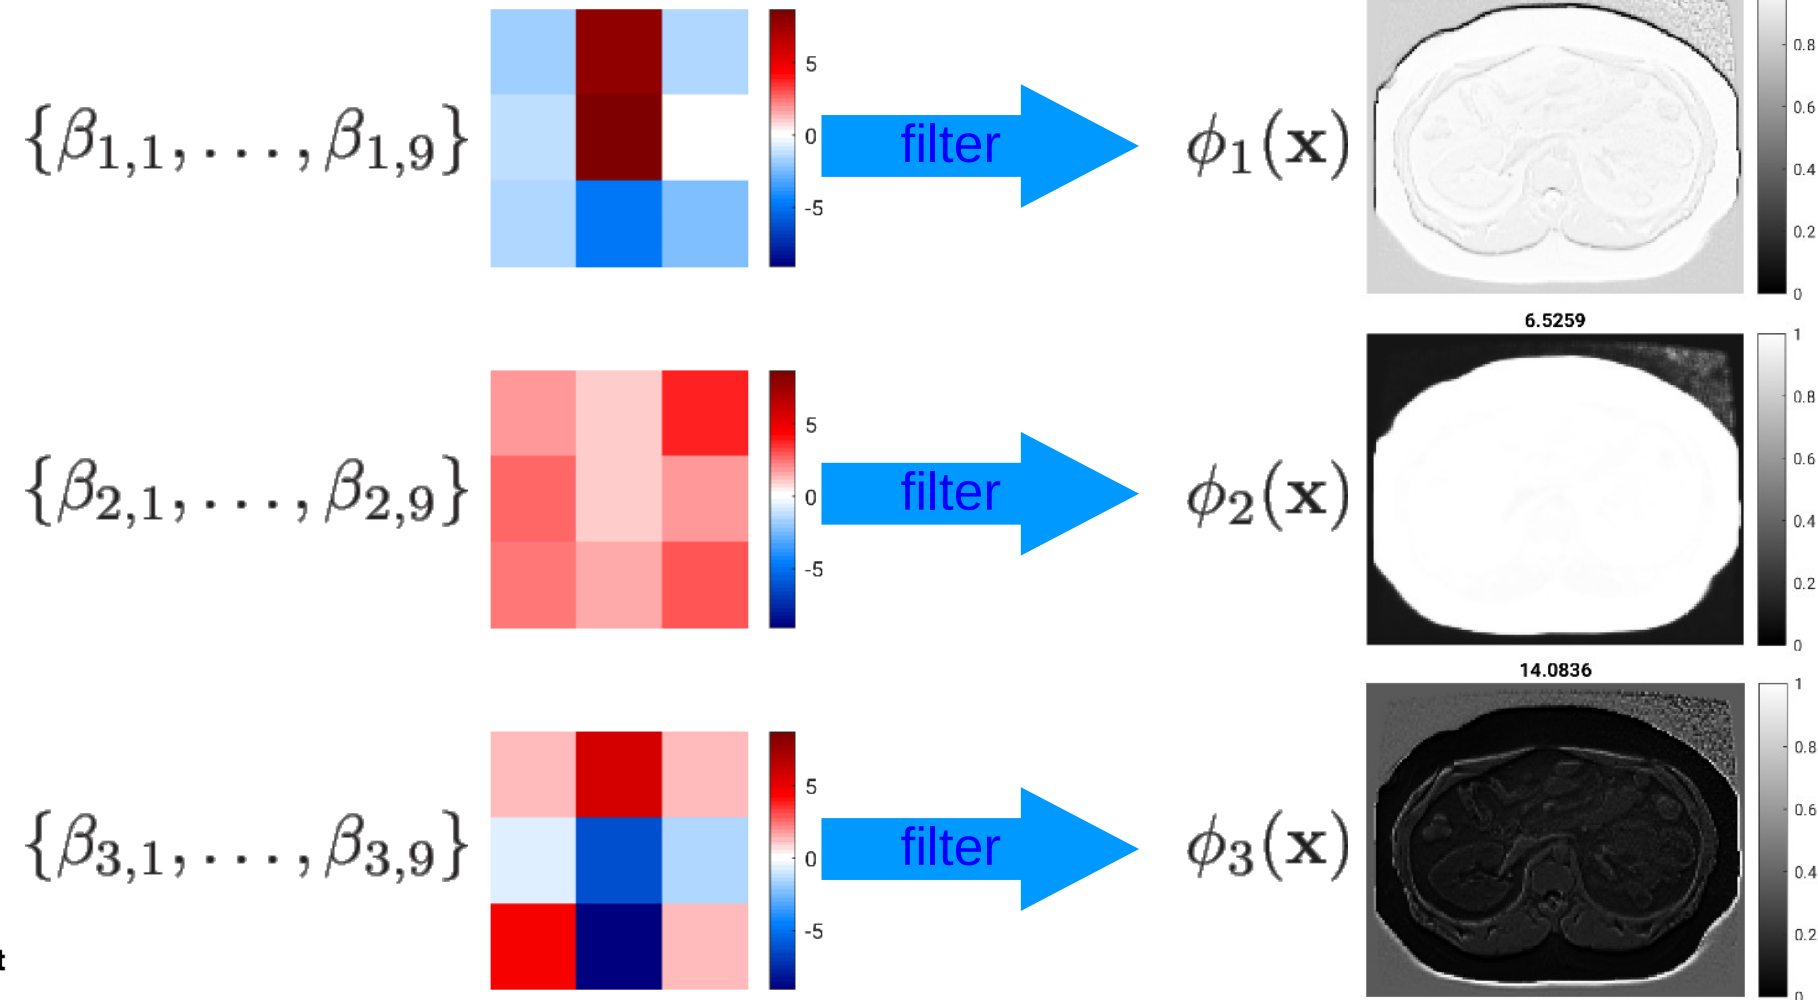
Applying
the trained
classifier
on new
data:



Applying
the trained
classifier
on new
data:



Filtering operations can
be implemented using
convolutions
=> “convolutional neural
network”



Neural networks = ultimate solution?

No model, only training data:



- No domain expertise needed
- Very easy to train and deploy
- Super fast (GPUs)



- Training data often very hard to get in medical imaging!
- Scanning hardware/software/protocol changes routinely!